Yelp Business Data Analysis: Feature Selection, Ratings Prediction, Sentiment Analysis

Nikhil Patil

Problem Description

- Feature Selection
- Ratings Prediction
- Sentiment Analysis
- Geospatial Visualization of Ratings

Dataset Description:

- Sub-datasets: business, user review, check-in. Total size: 5 GB.
- Business: 144K businesses with stars, attributes, categories
- E.g. parking availability, happy hour, drive thru, restaurants table service etc.
- Review: 4 M user comments, useful votes
- Check-in: check-in counts
- Challenges:
  - Data extraction
  - Integration of sub-datasets
  - Categorical features: quantification of True or False

Feature Selection

- Map Reduce Jobs.
- Frequently occurred attributes, categories in high star rated businesses: Suggestion to new businesses.
- User review count, available attributes count, available categories count, total check in counts and stars for each business.
- Machine learning approach to predict stars, based on these features.

BusinessAcceptsCreditCards 103200; BikeParking 51914
BusinessParking 49767; RestaurantsTakeOut 45544
GoodForKids 44115; RestaurantsGoodForGroups 40887
WheelchairAccessible 34799; GoodForMeal 25093
RestaurantsTableService 24423; Ambience 21965
HasTV 21233; OutdoorSeating 19510
RestaurantsReservations 17378; Caters 14938
ByAppointmentOnly 14201; RestaurantsDelivery 10162
BestNights 5632; HappyHour 5473
AcceptsInsurance 5255; Music 4850
DogsAllowed 3089; DriveThru 2154
GoodForDancing 1940; HairSpecializesIn 1079
CoatCheck 1011; RestaurantsCounterService 246
BusinessAcceptsBitcoin 178; DietaryRestrictions 155
Corkage 140; BYOB 47
Open24Hours 29}

Architecture: Feature Selection & Ratings Prediction

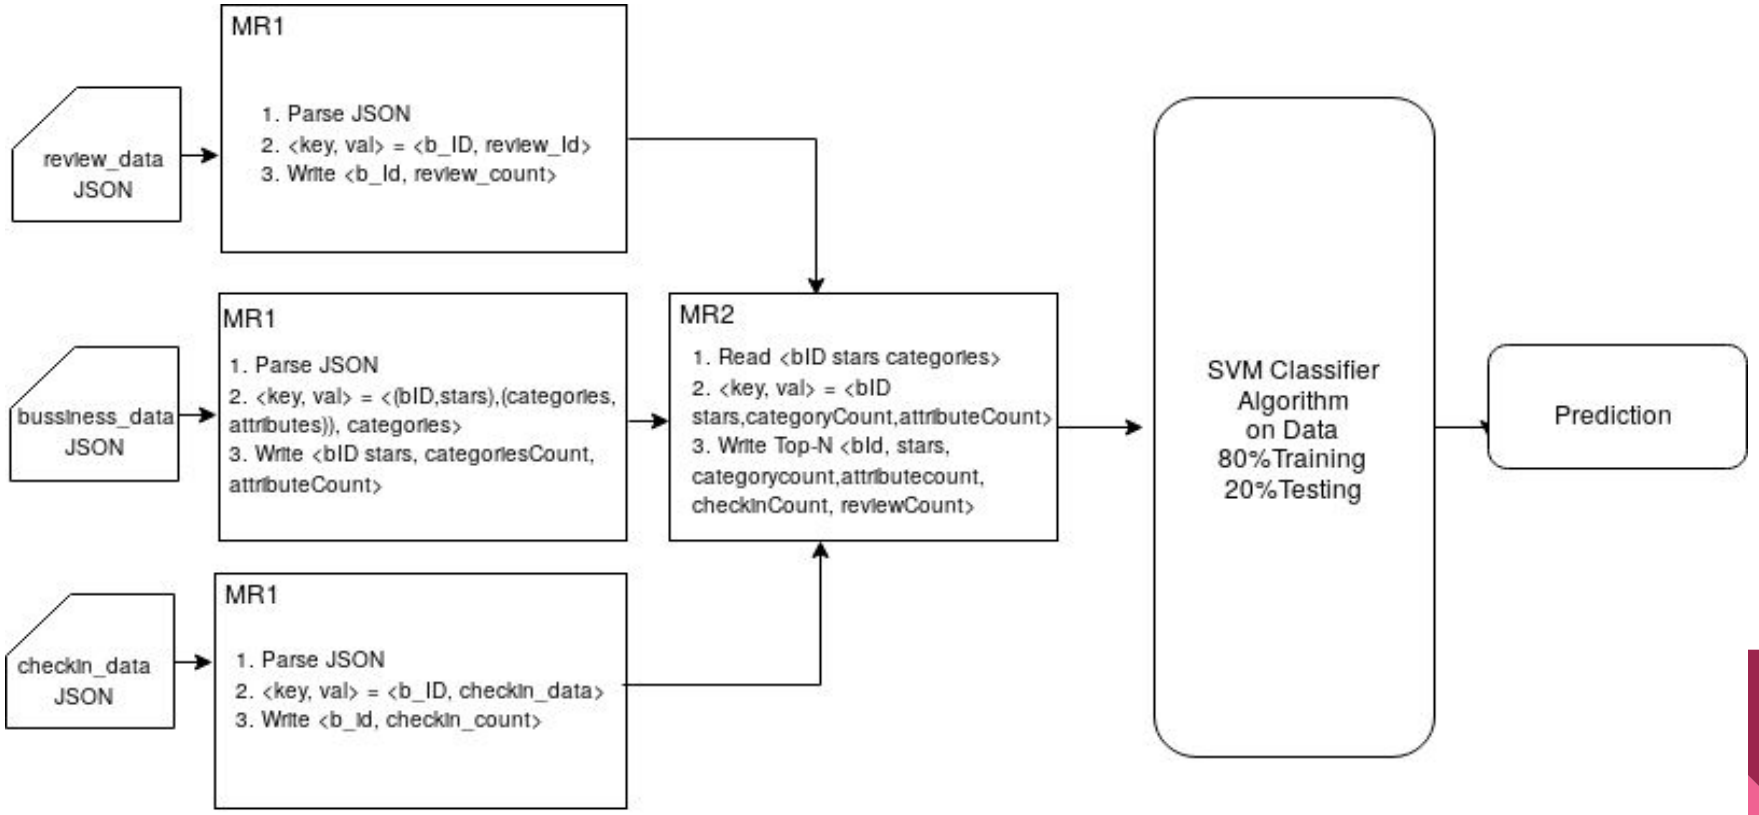

Ratings Prediction using Classification Algorithms

- Feature Spaces (X_train, test):
  - Count-based (4): [business_id, review count, attribute count, category count, checkin count]
  - Categorical (1 or 0) (31): [BusinessParking, HasTV, BYOB, etc.]
- Labels (y_train, test): Star ratings. [0,0.5,1,1.5,2,2.5,3,3.5,4,4.5,5] -> [0,1,2,3,4,5,6,7,8,9,10]
- Training Set: 75% (~110K); Test Set: 25% (~30K)
- Algorithms:
  - Regression: linear regression
  - Classification: SVM_linear, SVM_rbf, SGD

Evaluation of Analysis:

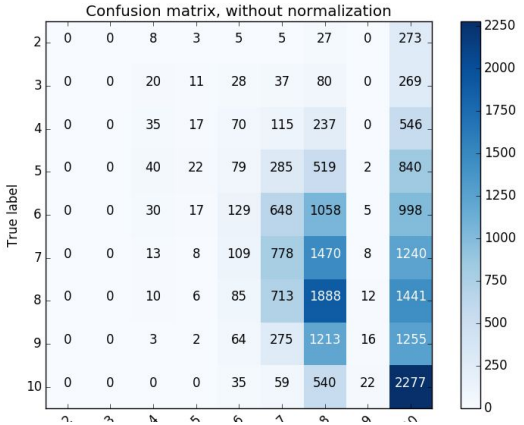

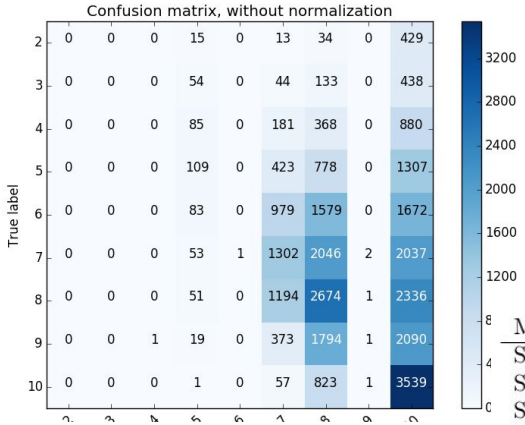

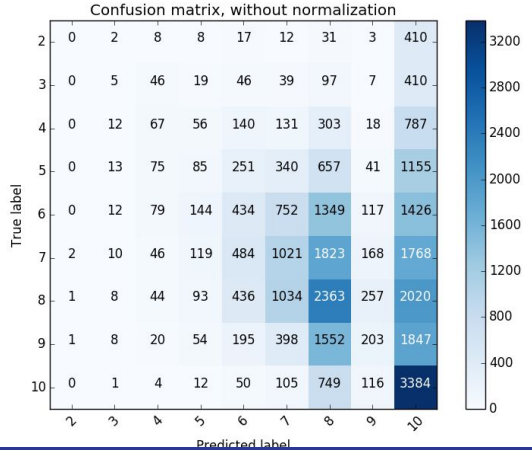

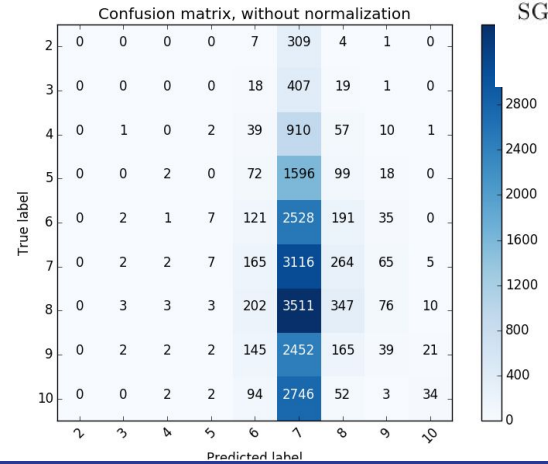

- Precision = TP/(TP+FP)
- Recall = TP/(TP+FN)
- F1 = 2*P*R/(P+R)

Model	Mean Precision	Mean Recall	Mean F1 score
SVM_linear (C=1)	0.24	0.26	0.19
SVM_rbf (C=1,g=1/31)	0.19	0.25	0.18
SVM_rbf (C=1000,g=0.3)	0.23	0.25	0.20
SGD	0.21	0.20	0.12

Table 3: Precision and Recall score for the models

Sentiment Analysis of Reviews on Business

- Sentiment score for each review
- Sentiment score = (n_Positive - n_Negative)/(n_Positive + n_Negative)
- Priority based sorting of users - stars, sentiment score, useful votes
- Top-N user suggestion.
- Average sentiment score for each business.
- Linear regression based star prediction using average sentiment score of business.
- Mean Square Error = 0.23

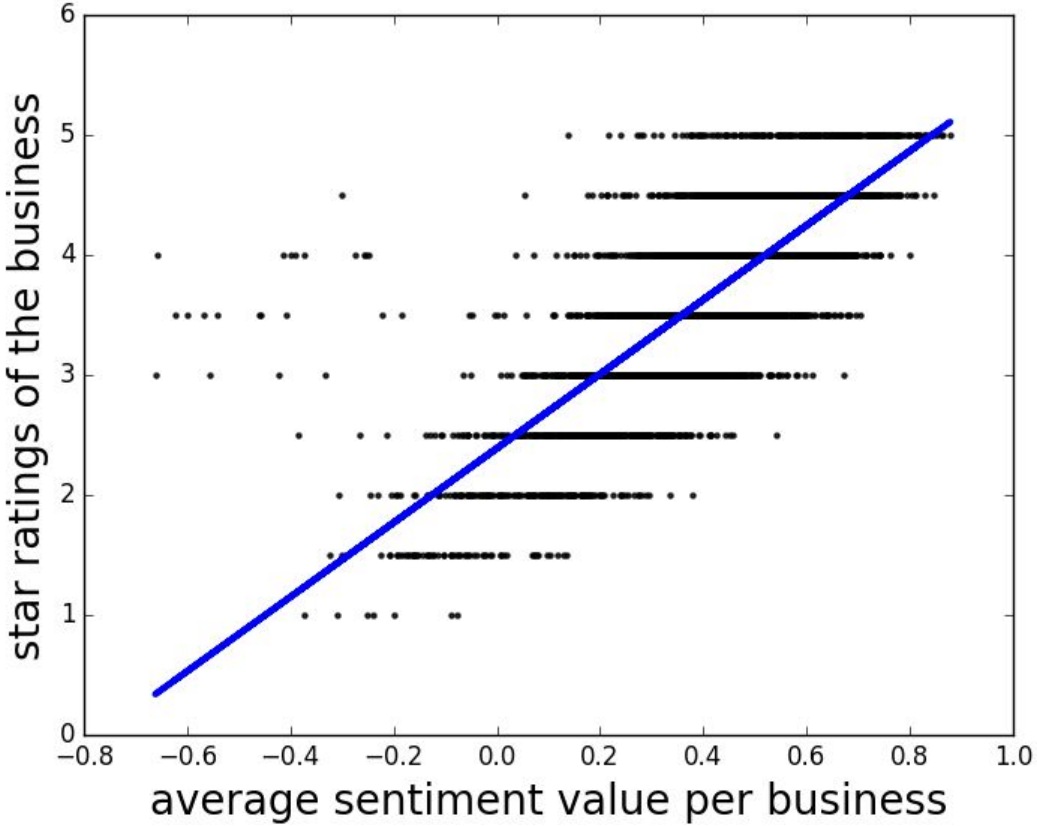

BusinessId	Stars	Sentiment Value
-2ToCaDFpTNmmg3QFzxcWg	1.5	-0.047114883
--DaPTJW3-tB1vP-PfdTEg	3.5	0.55616015
-AtzcXIwEP6yO7rM9CM9ww	4.0	0.54652196
-AxKgZHxyV-oBBHNyOESAg	5.0	0.66816974
HaRN97OjSnUJlHk21QAIXw	1.5	-0.0071489513
-B8uga7IGEQijKERiwuz7A	3.5	0.42080826
3quHwBY8XLalsx9m8e6Xuw	5.0	0.60437346
3rJJeMOkVi_wEura5UCQqw	3.0	0.38342175

Geospatial Visualization of Ratings

- Visualization to reveal trends in star ratings within a city
- Parse business_data in MapReduce with ‘city’ as key
- Get all Business_id, Latitude, Longitude, and Stars
- Use Google fusion table API for heatmap

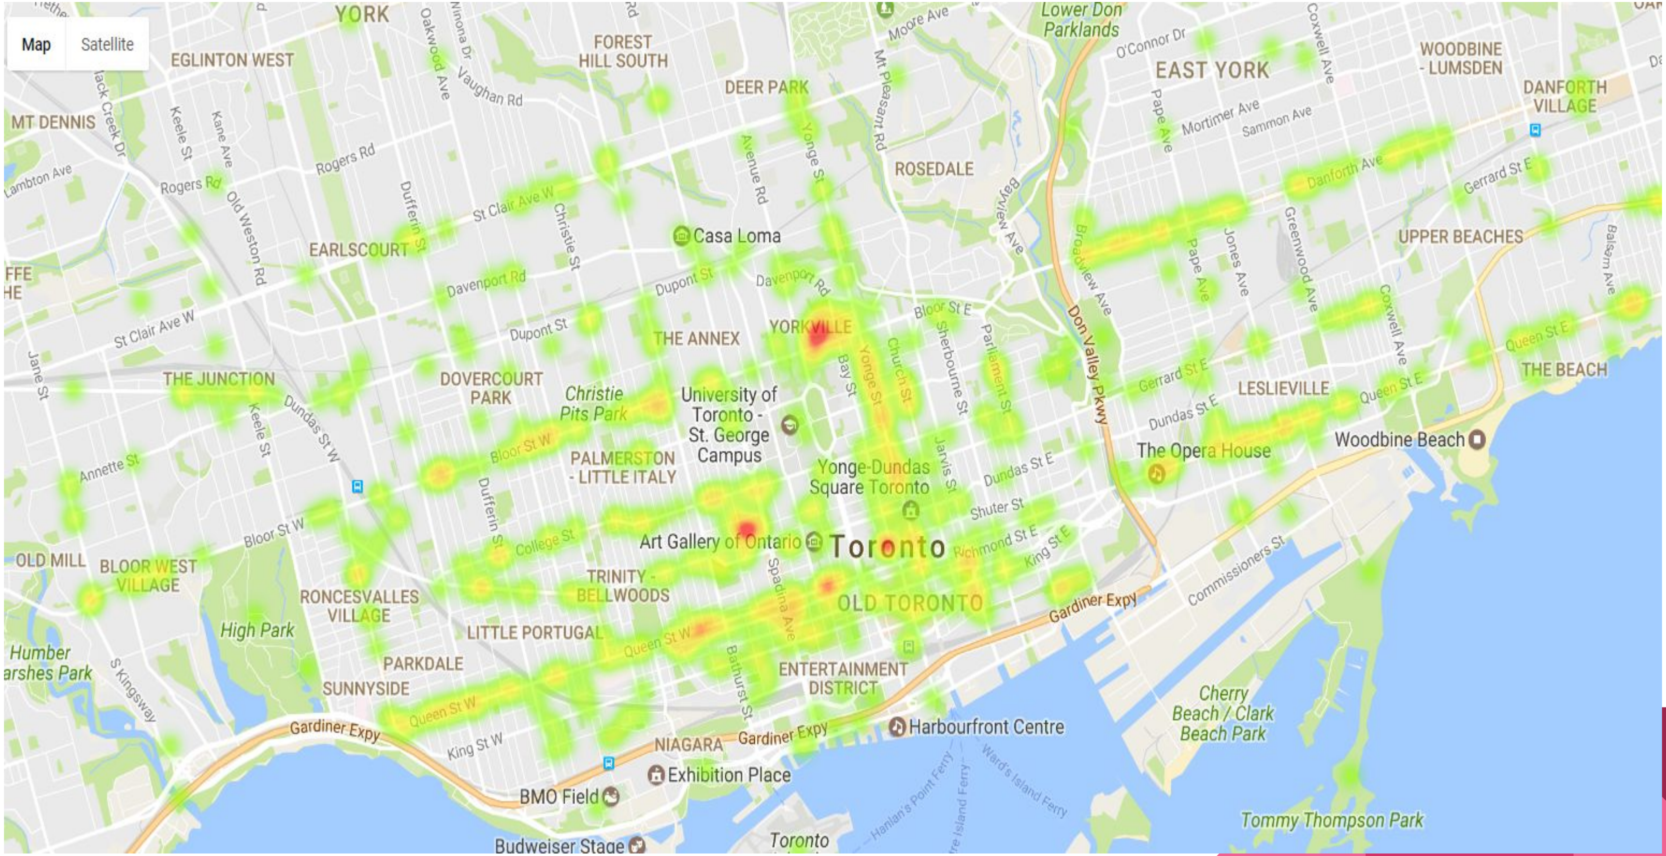

Map
Satellite
YORK
EGLINTON WEST
FOREST HILL SOUTH
DEER PARK
Lower Don Parklands
EAST YORK
WOODBINE - LUMSDEN
DANFORTH VILLAGE
MT DENNIS
ROSEDALE
EARLSCOURT
Casa Loma
UPPER BEACHES
YORKVILLE
THE ANNEX
THE JUNCTION
DOVERCOURT PARK
Christie Pits Park
University of Toronto - St. George Campus
LESLIEVILLE
THE BEACH
PALMERSTON - LITTLE ITALY
Yonge-Dundas Square Toronto
The Opera House
Woodbine Beach
Art Gallery of Ontario
Toronto
OLD MILL
BLOOR WEST VILLAGE
TRINITY - BELLWOODS
RONCESVALLES VILLAGE
OLD TORONTO
High Park
LITTLE PORTUGAL
Humber Marshes Park
PARKDALE
SUNNYSIDE
ENTERTAINMENT DISTRICT
Cherry Beach / Clark Beach Park
Harbourfront Centre
NIAGARA
Exhibition Place
BMO Field
Tommy Thompson Park
Budweiser Stage
Gardiner Expy
Don Valley Pkwy

Conclusion

- We did frequency based feature selection for attributes of business.
- Ratings prediction formulated as a 11-classification problem.
- To improve predictions:
  - Need better features
  - Need to optimize parameters for classification algorithms
- Sentiment Analysis: Unigrams-based sentiment results in positive correlation between sentiment value and the star ratings.
- Visualization helps to decide what areas, streets in a city have higher density of 4.5 or 5 stars.

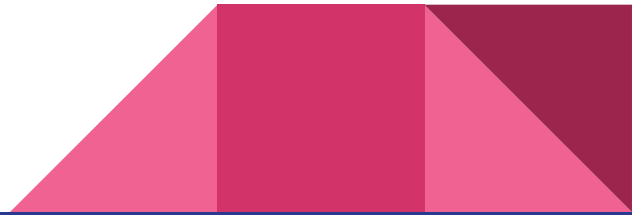

Questions?

Thank You.